

# MCMC techniques

Prosjekt 1 in MA8702. Written by Elling Svee.

January 23, 2026

## 1. Metropolis-Hastings for bivariate densities

In this project we implement three variations of the Metropolis-Hastings (MH). For evaluating the algorithms we consider three bivariate target densities for  $\mathbf{x} = (x, y)^T$ :

**Gaussian distribution** The first is a bivariate Gaussian distribution with correlation. Its probability density function (PDF) is

$$\pi(\mathbf{x}) = \frac{1}{2\pi\det(\Sigma)^{\frac{1}{2}}} \exp\left(-\frac{1}{2}\mathbf{x}^T\Sigma^{-1}\mathbf{x}\right) \quad (1)$$

where  $\Sigma$  has 1 on the diagonal and 0.9 on the off diagonals.

**Multimodal density** The second is a multimodal density constructed as a mixture of Gaussian densities. Its PDF is

$$\pi(\mathbf{x}) = \sum_{i=1}^3 w_i \frac{1}{2\pi\det(\Sigma_i)^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_i)^T\Sigma_i^{-1}(\mathbf{x} - \boldsymbol{\mu}_i)\right), \quad (2)$$

with  $w_i = 1/3$  for  $i = 1, 2, 3$ . The means are  $\boldsymbol{\mu}_1 = (-1.5, -1.5)^T$ ,  $\boldsymbol{\mu}_2 = (1.5, 1.5)^T$  and  $\boldsymbol{\mu}_3 = (-2, 2)^T$ , and the covariance matrices all have correlation 0 and variances  $\sigma_1^2 = \sigma_2^2 = 1$  and  $\sigma_3^2 = 0.8$ .

**Volcano density** Lastly we consider a volcano-shaped density with PDF

$$\pi(\mathbf{x}) \propto \frac{1}{2\pi} \exp\left(-\frac{1}{2}\mathbf{x}^T\mathbf{x}\right)(\mathbf{x}^T\mathbf{x} + 0.25) \quad (3)$$

Figure 1 visualize the three densities on a grid covering  $[-5, 5] \times [-5, 5]$ . The grid spacing is 0.1, which gives in total  $101 \times 101$  grid cells.




Figure 1: Bivariate densities

## 2. Theoretical background for Metropolis-Hastings

TODO

## 3. Code setup

The three MCMC algorithms are implemented in Python using the JAX (Bradbury et al., 2018) library for linear algebra and automatic differentiation (AD). Outside of the main implementation-scripts, we use the following files:

`densities.py` Implementation of the three densities.

```
import jax
import jax.numpy as jnp
import jax.scipy.stats as stats

def logdensity_mvnm(x):
    cov = jnp.array([[1.0, 0.8], [0.8, 1.0]])
    return stats.multivariate_normal.logpdf(x, mean=jnp.zeros(2), cov=cov)

def logdensity_multimodal(x):
    w = jnp.ones((3,)) / 3.0
    means = jnp.array([-1.5, -1.5], [1.5, 1.5], [-2.0, 2.0])

    # All three covariance matrices with 0 on the off-diagonal elements.
    # The first two have 1.0 on the diagonal, the last has 0.8.
    covs = jnp.array([jnp.eye(2), jnp.eye(2), 0.8 * jnp.eye(2)])

    # Compute log density for each component
    log_components = jax.vmap(
        lambda mean, cov: stats.multivariate_normal.logpdf(x, mean=mean, cov=cov),
        in_axes=(0, 0),
    )(means, covs)

    # Log-sum-exp trick: log(sum w_i * exp(log_p_i)) = logsumexp(log(w_i) +
    log_p_i)
    log_w = jnp.log(w)
    return jax.scipy.special.logsumexp(log_w + log_components)

def logdensity_volcano(x):
    xtx = jnp.sum(x**2)
    norm_const = 1.0 / (2 * jnp.pi)
    return jnp.log(norm_const) + jnp.log(xtx + 0.25) - 0.5 * xtx
```

`inference.py` For running the chains.

```
import jax

def inference_loop(rng_key, kernel, initial_state, num_samples):
    @jax.jit
    def one_step(state, rng_key):
        new_state, info = kernel(rng_key, state)
        return new_state, (new_state, info)

    keys = jax.random.split(rng_key, num_samples)
    _, (states, infos) = jax.lax.scan(one_step, initial_state, keys)

    return states, infos
```

`autocorr.py` Computing the autocorrelations.

```
import jax.numpy as jnp
```

```

def autocorr(x, max_lag=100, normalize=True):
    x = jnp.asarray(x)
    x = x - jnp.mean(x)

    n = x.shape[0]

    if max_lag is None:
        max_lag = n - 1
    max_lag = jnp.minimum(max_lag, n - 1)

    # FFT with zero-padding
    f = jnp.fft.fft(x, n=2 * n)
    ac = jnp.fft.ifft(f * jnp.conj(f)).real

    # keep only requested lags
    ac = ac[: max_lag + 1]

    if normalize:
        ac = ac / ac[0]

    return ac

```

tuning\_experiment.py Running the algorithms for different  $\sigma$ -values, and plotting results.

```

from os import PathLike
from typing import Callable

import jax
import jax.numpy as jnp
import matplotlib.pyplot as plt

from scripts.autocorr import autocorr
from scripts.inference import inference_loop


def run_tuning_experiment(
    init_fn: Callable,
    build_kernel_fn: Callable,
    logdensity_fn: Callable,
    filename: PathLike[str],
    sigma_values=[0.1, 0.5, 1.0, 1.5],
) -> None:
    """Run RWMH tuning experiment with different proposal stddev (sigma) values
    and plot the results.
    """
    # Run with different sigma values
    num_steps = 10000
    burnin = 1000

    key = jax.random.key(42)
    initial_pos = jnp.array([0.0, 0.0])

    fig, axes = plt.subplots(len(sigma_values), 3, figsize=(15, 4 *

```

```

len(sigma_values)))
fig_scatter, axes_scatter = plt.subplots(
    1, len(sigma_values), figsize=(5 * len(sigma_values), 5)
)

initial_state = init_fn(initial_pos, logdensity_fn)
for i, sigma in enumerate(sigma_values):
    kernel = build_kernel_fn(logdensity_fn, sigma)

    # Run chain
    key, subkey = jax.random.split(key)
    samples, infos = inference_loop(subkey, kernel, initial_state, num_steps)

    # Remove burnin
    positions = samples.position[burnin:]
    accept = infos.is_accepted[burnin:]

    # Acceptance rate
    acc_rate = float(jnp.mean(accept))

    row_title = rf"$\sigma = {sigma}$, Accept rate: {acc_rate:.3f}"

    # Trace plots
    axes[i, 0].plot(positions[:, 0], linewidth=0.5)
    axes[i, 0].set_ylabel("Trace x")

    axes[i, 1].plot(positions[:, 1], linewidth=0.5)
    axes[i, 1].set_ylabel("Trace y")
    axes[i, 1].set_title(row_title)

    # Autocorrelation
    acf = autocorr(positions[:, 0])
    axes[i, 2].bar(range(len(acf)), acf, width=1.0)
    axes[i, 2].set_ylabel("ACF")
    axes[i, 2].set_ylim([-0.2, 1.0])

    # 2D Scatter plot
    axes_scatter[i].scatter(positions[:, 0], positions[:, 1], alpha=0.5)
    axes_scatter[i].set_xlabel("x")
    axes_scatter[i].set_ylabel("y")
    axes_scatter[i].set_title(row_title)
    axes_scatter[i].axis("equal")
    axes_scatter[i].grid()

# Save the two figure separately
fig.tight_layout()
fig.savefig(f"{filename}.svg")
plt.close(fig)

fig_scatter.tight_layout()
fig_scatter.savefig(f"{filename}_scatter.svg")
plt.close(fig_scatter)

```

## 4. Random-walk Metropolis-Hastings

The first algorithm we implement is the Random-walk MH.

### 4.1. Theory

TODO

### 4.2. Implementation

`random_walk.py` Building the Random-walk MH kernel

```
import jax
from jax import Array
import jax.numpy as jnp
from typing import NamedTuple, Callable

class RWState(NamedTuple):
    position: Array
    logdensity: Array

class RWInfo(NamedTuple):
    acceptance_rate: Array
    is_accepted: Array
    proposal: RWState

def init(position: Array, logdensity_fn: Callable) -> RWState:
    return RWState(position, logdensity_fn(position))

def build_kernel(logdensity_fn: Callable, step_size: float) -> Callable:
    """Build a Random Walk Rosenbluth-Metropolis-Hastings kernel

    Returns
    ------
    A kernel that takes a rng_key and a Pytree that contains the current state
    of the chain and that returns a new state of the chain along with
    information about the transition.
    """

    def kernel(
        rng_key: Array,
        state: RWState,
    ) -> tuple[RWState, RWInfo]:
        # Generate proposal: x' = x + step_size * N(0, I)
        key_proposal, key_accept = jax.random.split(rng_key)
        proposal = state.position + step_size * jax.random.normal(
            key_proposal, shape=state.position.shape
        )

        # Compute log probability at proposal
        proposal_logdensity = logdensity_fn(proposal)
```

```

# Compute acceptance ratio (symmetric proposal cancels out)
log_ratio = proposal_logdensity - state.logdensity
acceptance_prob = jnp.minimum(1.0, jnp.exp(log_ratio))

# Accept or reject
u = jax.random.uniform(key_accept)
accepted = u < acceptance_prob

# Update state (use lax.cond for cleaner scalar handling)
new_position = jax.lax.select(accepted, proposal, state.position)
new_logdensity = jax.lax.select(accepted, proposal_logdensity,
state.logdensity)
new_state = RWState(new_position, new_logdensity)

# Store info
info = RWInfo(acceptance_prob, accepted, RWState(proposal,
proposal_logdensity))

return new_state, info

return kernel

```

`rwmh_chain.py` Running the tuning experiment for the Random-walk MH.

```

import os
from pathlib import Path
from scripts.tuning_experiment import run_tuning_experiment
from scripts.random_walk import init, build_kernel
from scripts.densities import logdensity_multimodal, logdensity_mvnb,
logdensity_volcano

OUTPUT_DIR = Path("output/rwmh/")
os.makedirs(OUTPUT_DIR, exist_ok=True)

print("Running tuning experiment for Random Walk Metropolis-Hastings...")
print("- Multivariate Normal distribution...")
run_tuning_experiment(init, build_kernel, logdensity_mvnb, OUTPUT_DIR /
"tuning_mvnb")
print("- Multimodal distribution...")
run_tuning_experiment(
    init, build_kernel, logdensity_multimodal, OUTPUT_DIR / "tuning_multimodal"
)
print("- Volcano distribution...")
run_tuning_experiment(
    init, build_kernel, logdensity_volcano, OUTPUT_DIR / "tuning_volcano"
)

```

### 4.3. Gaussian distribution

Figure 2 shows the tuning experiment for the Gaussian distribution. Observe that  $\sigma = 1.5$  gives an acceptance rate of 0.264. This is closest to the theoretical optimal, and is therefore preferred.




Figure 2: Tuning experiment of Random-walk MH for Gaussian distribution

#### 4.4. Multimodal distribution

Figure 3 shows the tuning experiment for the multimodal distribution.  $\sigma = 1.5$  gives an acceptance rate of 0.503. This is closest to the theoretical optimal, and is therefore preferred.




Figure 3: Tuning experiment of Random-walk MH for multimodal distribution

#### 4.5. Volcano distribution

Figure 4 shows the tuning experiment for the volcano distribution.  $\sigma = 1.5$  gives an acceptance rate of 0.523. This is closest to the theoretical optimal, and is therefore preferred.




Figure 4: Tuning experiment of Random-walk MH for volcano distribution

## 5. Langevin Metropolis-Hastings

The second algorithm we implement is the Langevin MH.

### 5.1. Theory

TODO

The scaling limit literature indicates that the optimal acceptance probability is approximately 0.57 (Dunson and Johndrow, 2020).

### 5.2. Implementation

`langevin.py` Building the Langevin MH kernel

```
from typing import Callable, NamedTuple
```

```

import jax
import jax.numpy as jnp
from jax import Array

class LangevinState(NamedTuple):
    position: Array
    logdensity: Array
    logdensity_grad: Array

class LangevinInfo(NamedTuple):
    acceptance_rate: Array
    is_accepted: Array
    proposal: LangevinState

def init(position: Array, logdensity_fn: Callable) -> LangevinState:
    logdensity, logdensity_grad = jax.value_and_grad(logdensity_fn)(position)
    return LangevinState(position, logdensity, logdensity_grad)

def build_kernel(logdensity_fn: Callable, step_size: float) -> Callable:
    def kernel(
        rng_key: Array,
        state: LangevinState,
    ) -> tuple[LangevinState, LangevinInfo]:
        # Generate proposal:  $x' = x + \text{step\_size} * N(0, I)$ 
        key_proposal, key_accept = jax.random.split(rng_key)

        proposal = (
            state.position
            + step_size * state.logdensity_grad
            + jnp.sqrt(2 * step_size)
            * jax.random.normal(key_proposal, shape=state.position.shape)
        )

        proposal_logdensity, proposal_logdensity_grad = jax.value_and_grad(
            logdensity_fn
        )(proposal)

        # Compute acceptance ratio (symmetric proposal cancels out)
        log_ratio = proposal_logdensity - state.logdensity

        # Compute the proposal densities  $q(x'|x)$  and  $q(x|x')$ 
        def log_proposal_density(from_pos, to_pos, from_grad):
            diff = to_pos - from_pos - step_size * from_grad
            return -0.5 * jnp.sum(diff**2) / (2 * step_size)

        log_q_forward = log_proposal_density(
            state.position, proposal, state.logdensity_grad
        )
        log_q_backward = log_proposal_density(
            proposal, state.position, proposal_logdensity_grad
        )

```

```

        )
    log_ratio += log_q_backward - log_q_forward
    acceptance_prob = jnp.minimum(1.0, jnp.exp(log_ratio))

    # Accept or reject
    u = jax.random.uniform(key_accept)
    accepted = u < acceptance_prob

    # Update state (use lax.cond for cleaner scalar handling)
    new_position = jax.lax.select(accepted, proposal, state.position)
    new_logdensity = jax.lax.select(accepted, proposal_logdensity,
state.logdensity)
    new_logdensity_grad = jax.lax.select(
        accepted, proposal_logdensity_grad, state.logdensity_grad
    )
    new_state = LangevinState(new_position, new_logdensity,
new_logdensity_grad)

    # Store info
    info = LangevinInfo(
        acceptance_prob,
        accepted,
        LangevinState(proposal, proposal_logdensity,
proposal_logdensity_grad),
    )

    return new_state, info

return kernel

```

`langevin_chain.py` Running the tuning experiment for the Langevin MH.

```

import os
from pathlib import Path
from scripts.tuning_experiment import run_tuning_experiment
from scripts.langevin import init, build_kernel
from scripts.densities import logdensity_multimodal, logdensity_mvnb,
logdensity_volcano

OUTPUT_DIR = Path("output/langevin/")
os.makedirs(OUTPUT_DIR, exist_ok=True)

print("Running tuning experiment for Langevin Monte Carlo...")
print("- Multivariate Normal distribution...")
run_tuning_experiment(init, build_kernel, logdensity_mvnb, OUTPUT_DIR /
"tuning_mvnb")
print("- Multimodal distribution...")
run_tuning_experiment(
    init, build_kernel, logdensity_multimodal, OUTPUT_DIR / "tuning_multimodal"
)
print("- Volcano distribution...")
run_tuning_experiment(

```

```

    init, build_kernel, logdensity_volcano, OUTPUT_DIR / "tuning_volcano"
)

```

### 5.3. Gaussian distribution

Figure 5 shows the tuning experiment for the Gaussian distribution.  $\sigma = 0.5$  is closest to the optimal acceptance probability, and is therefore preferred.




Figure 5: Tuning experiment of Langevin MH for Gaussian distribution

### 5.4. Multimodal distribution

Figure 6 shows the tuning experiment for the multimodal distribution.  $\sigma = 1.0$  is closest to the optimal acceptance probability, and is therefore preferred.




Figure 6: Tuning experiment of Langevin MH for multimodal distribution

## 5.5. Volcano distribution

Figure 7 shows the tuning experiment for the volcano distribution.  $\sigma = 1.5$  is closest to the optimal acceptance probability, and is therefore preferred.




Figure 7: Tuning experiment of Langevin MH for multimodal distribution

## 6. Hamiltonian Metropolis-Hastings

The third algorithm we implement is the Hamiltonian MH.

### 6.1. Theory

TODO

### 6.2. Implementation

`hamiltonian.py` Building the Hamiltonian MH kernel

```
from typing import Callable, NamedTuple

import jax
import jax.numpy as jnp
```

```

from jax import Array

class HMCState(NamedTuple):
    position: Array
    logdensity: Array
    logdensity_grad: Array

class HMCInfo(NamedTuple):
    acceptance_rate: Array
    is_accepted: Array
    proposal: HMCState

def init(position: Array, logdensity_fn: Callable) -> HMCState:
    logdensity, logdensity_grad = jax.value_and_grad(logdensity_fn)(position)
    return HMCState(position, logdensity, logdensity_grad)

def hamiltonian(logdensity, momentum):
    kinetic = 0.5 * jnp.sum(momentum**2)
    potential = -logdensity
    return potential + kinetic

def leapfrog(
    position: Array,
    momentum: Array,
    logdensity_fn: Callable,
    step_size: float,
    num_steps: int,
):
    def body_fn(_, state):
        x, p, logp, grad = state

        # Half step momentum
        p = p + 0.5 * step_size * grad

        # Full step position
        x = x + step_size * p

        # Refresh gradient
        logp, grad = jax.value_and_grad(logdensity_fn)(x)

        # Half step momentum
        p = p + 0.5 * step_size * grad

    return x, p, logp, grad

log0, grad0 = jax.value_and_grad(logdensity_fn)(position)
position, momentum, logp, grad = jax.lax.fori_loop(
    0,
    num_steps,
)

```

```

        body_fn,
        (position, momentum, logp0, grad0),
    )

    return position, momentum, logp, grad

def build_kernel(
    logdensity_fn: Callable,
    step_size: float,
    num_steps: int = 10,
) -> Callable:
    def kernel(
        rng_key: Array,
        state: HMCState,
    ) -> tuple[HMCState, HMCInfo]:
        key_momentum, key_accept = jax.random.split(rng_key)

        # Sample momentum
        momentum0 = jax.random.normal(key_momentum, shape=state.position.shape)

        # Current Hamiltonian
        H = hamiltonian(state.logdensity, momentum0)

        # Propose new state via leapfrog integrator
        q_prop, p_prop, logp_prop, grad_prop = leapfrog(
            state.position,
            momentum0,
            logdensity_fn,
            step_size,
            num_steps,
        )

        # Proposed Hamiltonian
        H_prop = hamiltonian(logp_prop, p_prop)

        # Acceptance probability
        log_accept_ratio = H - H_prop
        acceptance_prob = jnp.minimum(1.0, jnp.exp(log_accept_ratio))

        # Accept or reject
        u = jax.random.uniform(key_accept)
        accepted = u < acceptance_prob

        new_state = HMCState(
            position=jax.lax.select(accepted, q_prop, state.position),
            logdensity=jax.lax.select(accepted, logp_prop, state.logdensity),
            logdensity_grad=jax.lax.select(accepted, grad_prop,
state.logdensity_grad),
        )

        info = HMCInfo(
            acceptance_rate=acceptance_prob,
            is_accepted=accepted,
            proposal=HMCState(q_prop, logp_prop, grad_prop),

```

```

    )

        return new_state, info

    return kernel

```

`hmc_chain.py` Running the tuning experiment for the Hamiltonian MH.

```

import os
from pathlib import Path

from scripts.densities import logdensity_multimodal, logdensity_mvnb,
logdensity_volcano
from scripts.hamiltonian import build_kernel, init
from scripts.tuning_experiment import run_tuning_experiment

OUTPUT_DIR = Path("output/hmc/")
os.makedirs(OUTPUT_DIR, exist_ok=True)

print("Running tuning experiment for Hamiltonian Monte Carlo...")

step_sizes = [0.01, 0.1, 0.25, 0.5]

print("- Multivariate Normal distribution...")
run_tuning_experiment(
    init,
    build_kernel,
    logdensity_mvnb,
    OUTPUT_DIR / "tuning_mvnb",
    sigma_values=step_sizes,
)
print("- Multimodal distribution...")
run_tuning_experiment(
    init,
    build_kernel,
    logdensity_multimodal,
    OUTPUT_DIR / "tuning_multimodal",
    sigma_values=step_sizes,
)
print("- Volcano distribution...")
run_tuning_experiment(
    init,
    build_kernel,
    logdensity_volcano,
    OUTPUT_DIR / "tuning_volcano",
    sigma_values=step_sizes,
)

```

### 6.3. Gaussian distribution

Figure 8 shows the tuning experiment for the Gaussian distribution. For Hamiltonian MC we prefer a high acceptance rate. Looking at the plots we see that  $\sigma = 0.5$  gives the smallest correlation between consecutive samples. As it still has a very high acceptance rate, this is preferred.




Figure 8: Tuning experiment of Hamiltonian MH for Gaussian distribution

#### 6.4. Multimodal distribution

Figure 9 shows the tuning experiment for the multimodal distribution. Again, the  $\sigma = 0.5$  gives fast-decreasing automatic while retaining a high acceptance rate. This is therefore preferred.




Figure 9: Tuning experiment of Hamiltonian MH for multimodal distribution

## 6.5. Volcano distribution

Figure 10 shows the tuning experiment for the volcano distribution. For the same reasons as for the previous other distributions, the  $\sigma = 0.5$  is preferred.




Figure 10: Tuning experiment of Hamiltonian MH for multimodal distribution

## Bibliography

- Bradbury, J., Frostig, R., Hawkins, P., Johnson, M.J., Leary, C., Maclaurin, D., Necula, G., Paszke, A., VanderPlas, J., Wanderman-Milne, S., Zhang, Q., 2018. JAX: composable transformations of Python+NumPy programs [WWW Document].. URL <http://github.com/jax-ml/jax>
- Dunson, D.B., Johndrow, J.E., 2020. The Hastings algorithm at fifty. *Biometrika* 107, 1–23.. <https://doi.org/10.1093/biomet/asz066>